

# AlphaFold2 on Alpine

RC Short Course

11 November 2024

Layla Freeborn



Research Computing  
Office of Information and Technology

UNIVERSITY OF COLORADO BOULDER

# Slides



<https://bit.ly/af2alpine>



Research Computing  
Office of Information and Technology

UNIVERSITY OF COLORADO BOULDER

# CURC User Support



## Meet the User Support Team



Layla  
Freeborn



Brandon  
Reyes



Andy  
Monaghan



Michael  
Schneider



John  
Reiland



Dylan  
Gottlieb



Mohal  
Khandelwal



Ragan  
Lee



Research Computing  
Office of Information and Technology

UNIVERSITY OF COLORADO BOULDER

# Background: The Protein Folding Problem

- What a protein does (it's biological function) depends on its 3D structure
- Figuring out the what shapes a protein folds into has been a grand challenge in biology for 50+ years

**Can we predict a protein's 3D structure based on its 1D aa sequence?**

This is computationally **VERY** intensive!



# Background

Article | [Open access](#) | Published: 15 July 2021

## Highly accurate protein structure prediction with AlphaFold

[John Jumper](#) , [Richard Evans](#), [Alexander Pritzel](#)

[Kathryn Tunyasuvunakool](#), [Russ Bates](#), [Augustin Žídek](#)

[Meyer](#), [Simon A. A. Kohl](#), [Andrew J. Ballard](#), [Andrea](#)

[Nikolov](#), [Rishabh Jain](#), [Jonas Adler](#), [Trevor Back](#), [Stijn](#)

[Zielinski](#), ... [Demis Hassabis](#)  [+ Show authors](#)

[Nature](#) **596**, 583–589 (2021) | [Cite this article](#)

**1.56m** Accesses | **8815** Citations | **3559** Altmetric | [Metrics](#)

### Access & Citations

**1.57m**

Article Accesses

**8815**

[Web of Science](#)

**1237**

[CrossRef](#)

### Online attention



5229 tweeters

318 news outlets

17 Video uploaders

104 blogs

7 Redditors

36 Wikipedia page

15600 Mendeley

6 Facebook pages

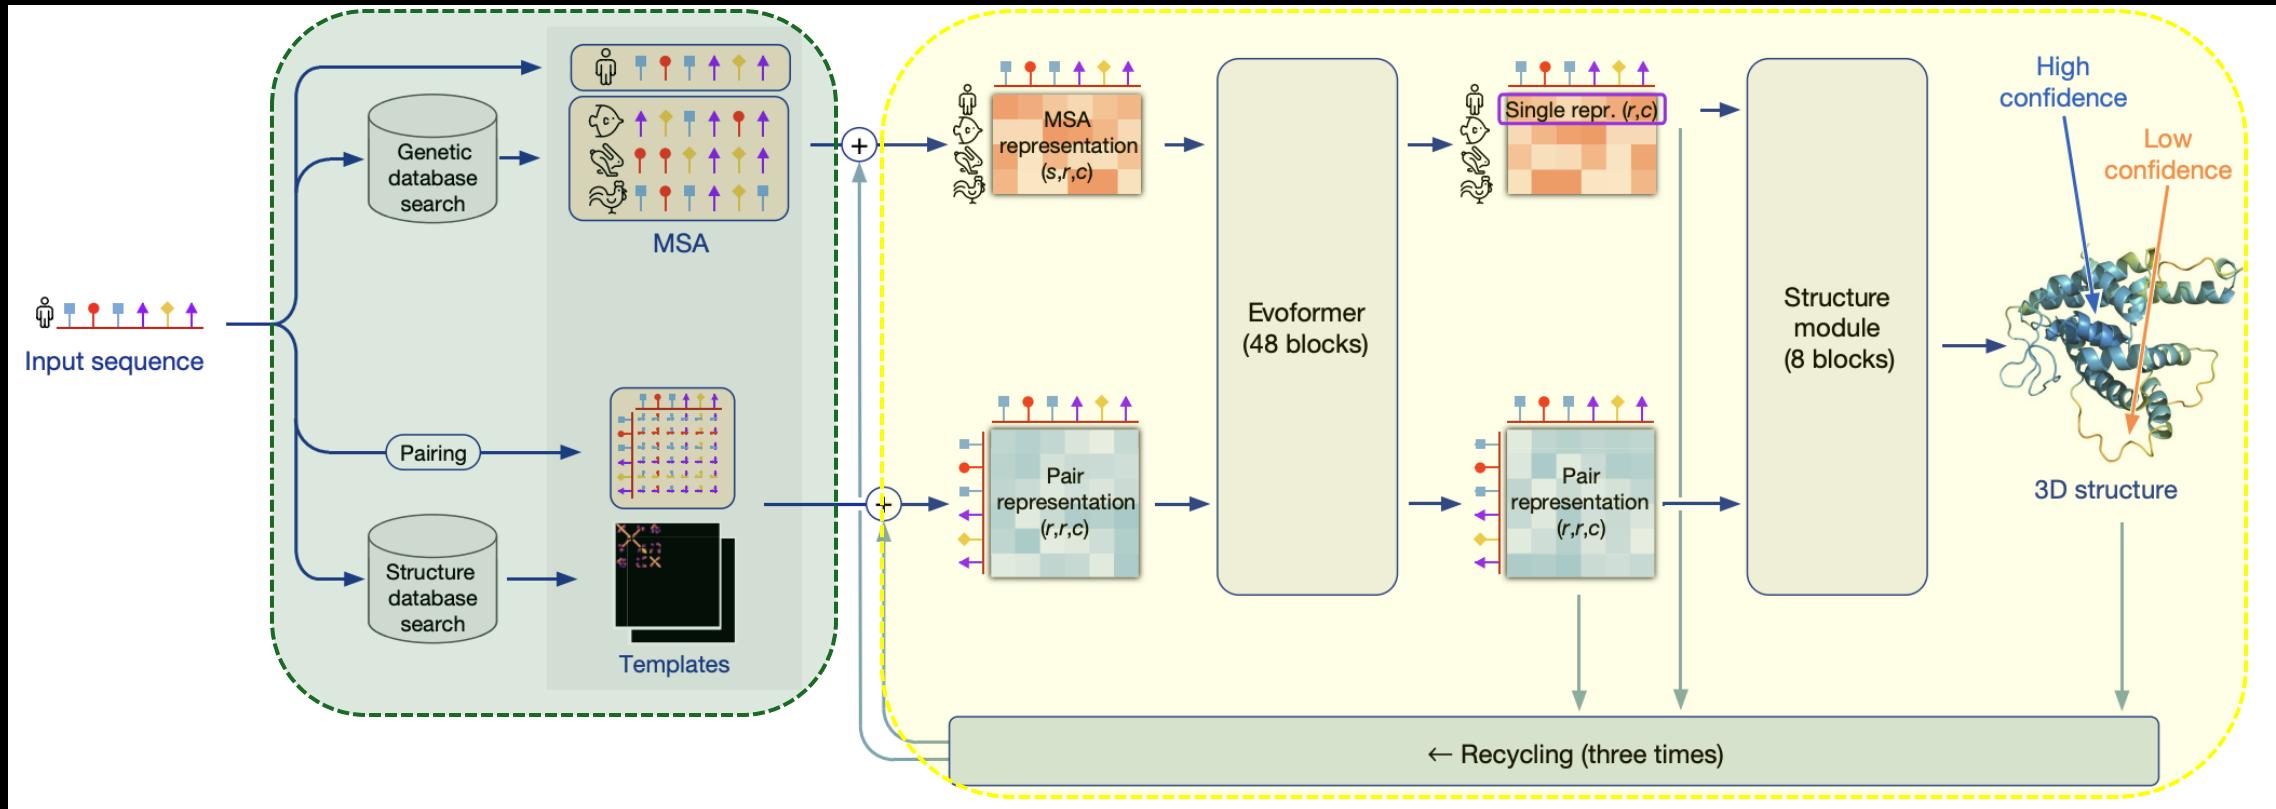
15 F1000

This article is in the 99<sup>th</sup> percentile (ranked 108<sup>th</sup>) of the 449,474 tracked articles of a similar age in all journals and the 98<sup>th</sup> percentile (ranked 15<sup>th</sup>) of the 927 tracked articles of a similar age in *Nature*

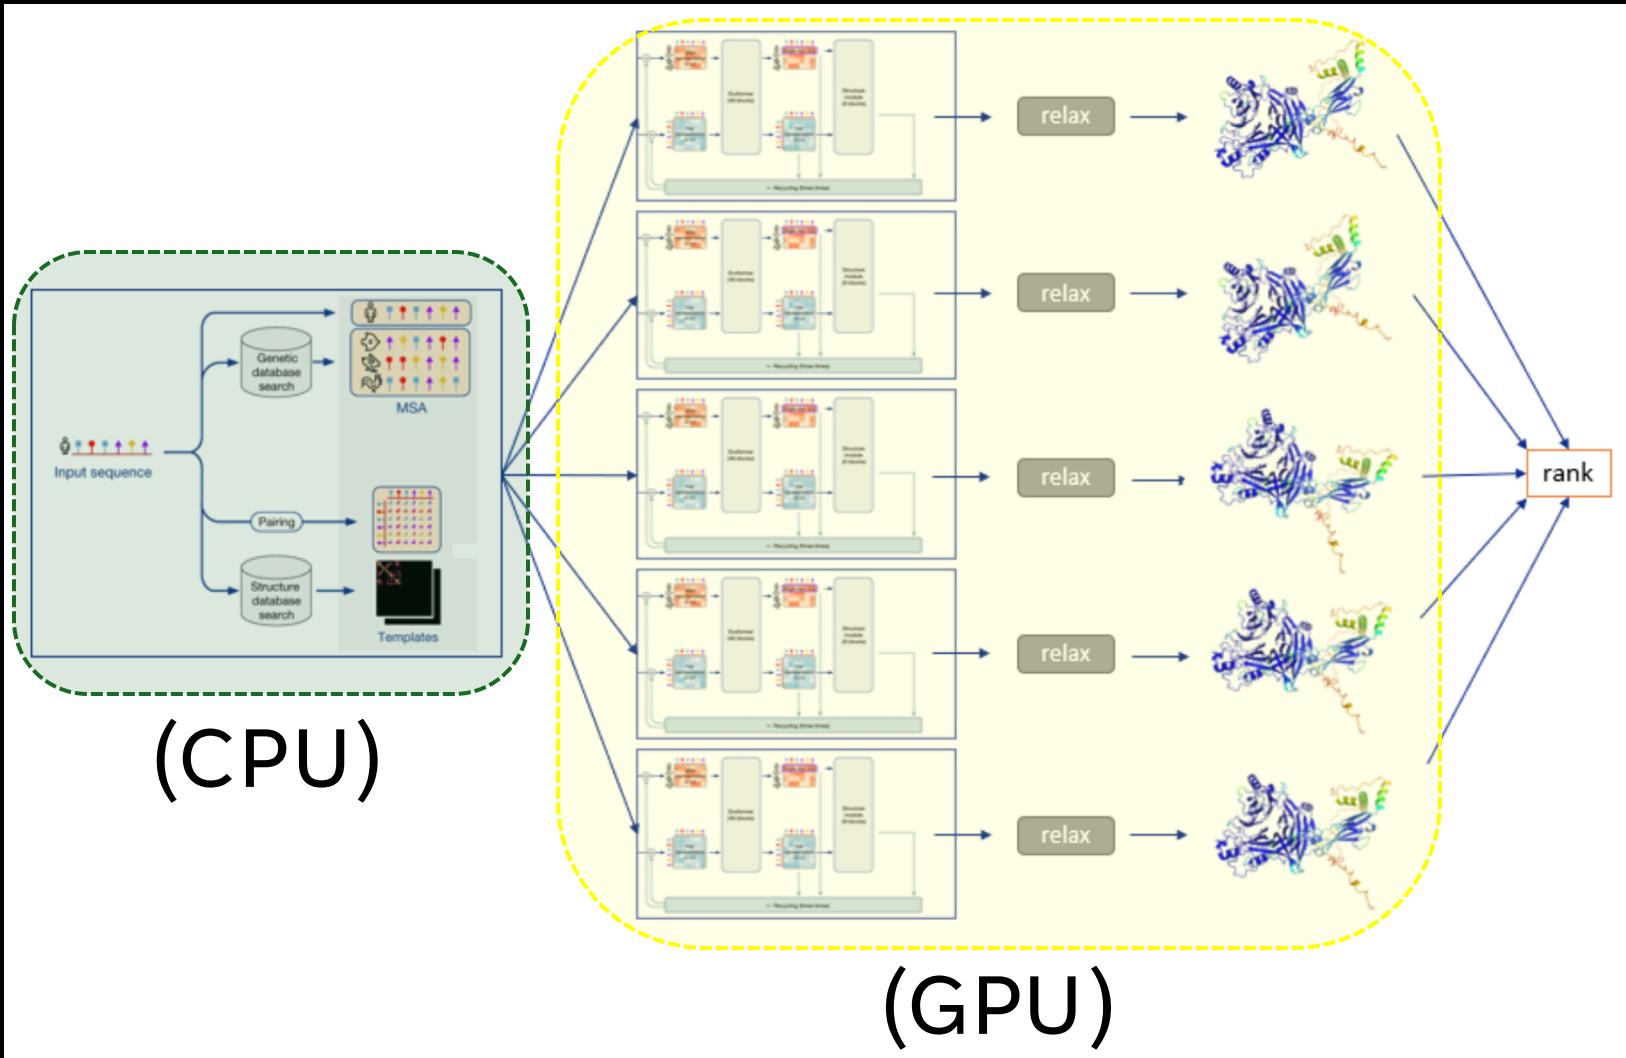
[View more on Altmetric](#)



# The Pipeline



# The Pipeline: CPU vs GPU



# AlphaFold Inputs

- FASTA file (e.g., `dummy.fasta`, `T1061.fa`)

```
>T1061 T5 phage tail subcomplex, T5 phage, 949 residues|  
MKKILDSAKNYLNTHDKLKTACLIALELPSSSGSAATYIYLTDYFRDVTYNGILYRSGKVKSISSHKQNRQLSIGSLSFTITGTAEDEV  
KLVQNGVSFLDRGITIHQAIINEEGNILPVDPTDGPLLFFRGRITGGGIKDNDVNNTSGIGTSVITWNCSNQFYDFDRVNGRYTDDASHRG  
LEVNVNTLQPSNGAKRPEYQEDYGFHSNKSTTILAKYQVKEERYKLQSKKKLFGLRSYSLKYYETVTKEVLDNFNLAAKFIPVVYGV  
QKIPGIPIFADTELNNPNIVVVYAFAEGEIDGFLDFYIGDSPMICFDETSDTRTCFGRKKIVGDTMHRLAAGTSTSQPSVHGQEYKYN  
DGNGDIRIWTFHGKPDQTAAQVLVDIAKKKGFYLNQNQNGNGPEYWDSRYKLLDTAYAIVRFTINENRTEIPEISAEVQGKKVKVYN  
SDGTIKADKTSLNGIWQLMDYLTSDRYGADITLDQFPLQKVISEAKILDIDIIDESYQTSWQPYWRYVGWN  
DPLSENQIVQLNTILDSESVFKNVQGILESFGGA  
INNLSGEYRITVEKYSTNPLRINF  
LDTYGDLDLS  
DTTGRNK  
FNSVQASL  
VDPAL  
SWKTN  
SITFYNS  
KFKEQDK  
GLDKKL  
QLSFAN  
ITNYYT  
TARSYAD  
RELKK  
SRY  
RTLS  
FSVPY  
KFIGIE  
PN  
DPIAF  
TYERY  
GW  
W  
K  
DK  
FF  
LV  
DE  
VENT  
RD  
G  
K  
IN  
L  
V  
L  
Q  
EY  
GED  
V  
FIN  
SE  
QVDNS  
GNDIP  
DISNN  
VLPP  
RDF  
KY  
TP  
PG  
VV  
GA  
IG  
K  
NG  
EL  
SW  
L  
PS  
LT  
NN  
VV  
YY  
SIAH  
SGH  
VNP  
YIV  
QQ  
LEN  
NP  
N  
ER  
MI  
QE  
EI  
IG  
P  
AG  
LA  
IF  
EL  
RA  
VD  
ING  
RR  
SP  
VT  
LS  
VD  
LNS  
AK  
N  
LS  
V  
VS  
N  
FR  
VV  
NT  
AS  
GD  
V  
TE  
F  
V  
GP  
DV  
K  
L  
AW  
DK  
I  
PE  
EE  
II  
I  
PE  
I  
YY  
T  
L  
E  
I  
Y  
D  
S  
Q  
D  
R  
M  
L  
R  
V  
R  
IED  
V  
Y  
T  
Y  
D  
Y  
L  
L  
T  
Y  
N  
K  
A  
D  
F  
A  
L  
L  
N  
S  
G  
A  
L  
G  
I  
N  
R  
K  
L  
F  
R  
I  
R  
A  
E  
G  
E  
N  
G  
E  
Q  
S  
V  
G  
W  
AT  
I
```

A FASTA file is the only *required* AlphaFold input.



# Alpine AlphaFold2 Module

Loading the AlphaFold module (`ml alphafold`)

- activates a conda environment containing all the necessary packages
- redirects `TMPDIR` to `/scratch/alpine/$USER`
- creates a shortcut to the AlphaFold run script
- sets convenient environment variables

`$CURC_AF_DBs`

`$CURC_AF_EXAMPLES`



# Alpine AlphaFold2 Module

```
-- This module loads AlphaFold

-- Set local variables
local PACKAGE_PREFIX = "/curc/sw/install/bio/alphafold/2.3.1/2.3.1_env"
local USER = os.getenv("USER")

-- Load dependencies
always_load("anaconda/2022.10")

-- Set paths
prepend_path("LD_LIBRARY_PATH", pathJoin(PACKAGE_PREFIX, "lib"))

-- Set environment variables
setenv("TMPDIR", (os.getenv("SCRATCHDIR")) .. "/" .. (os.getenv("USER")))
setenv("CURC_AF_DB", "/gpfs/alpine1/datasets/bioinformatics/alphafold")
setenv("CURC_AF_EXAMPLES", "/curc/sw/install/bio/alphafold/examples")

-- Create AlphaFold script shortcut
-- Note this will only work with bash; can add csh using $* (see https://lmod.readthedocs.io/en/latest/015_writing_modules.html#setting-aliases-and-shell-functions)
set_shell_function("run_alphaFold", 'bash /curc/sw/install/bio/alphafold/2.3.1/alphafold-2.3.1/run_alphaFold.sh "$@"')
```

sets environment variables/redirects  
TMPDIR

creates a shortcut to run\_alphaFold.sh



# Alpine AlphaFold2 Module

```
-- Activate alphafold conda environment
execute{cmd="conda activate /curc/sw/install/bio/alphafold/2.3.1/2.3.1_en
v",modeA={"load"}}
-- Software info:

help([[Alphafold:
AlphaFold is an artificial intelligence program developed by DeepMind, a
subsidiary of Alphabet, which performs predictions of protein structure.

For detailed instructions, go to:
https://github.com/deepmind/alphafold
]])
whatis("Version: 2.3.1")
whatis("Citing AlphaFold: Jumper, J et al. Highly accurate protein struct
ure prediction with AlphaFold. Nature (2021). Varadi, M et al. AlphaFold
Protein Structure Database: massively expanding the structural coverage o
f protein-sequence space with high-accuracy models. Nucleic Acids Researc
h (2022).")
```

activates the AlphaFold conda  
environment



# AlphaFold2 Required Parameters

```
(/curc/sw/install/bio/alphafold/2.3.1/2.3.1_env) [lafr9499@c3cpu-a2-u32-3 ~]$ run_alphaFold
```

Please make sure all required parameters are given

Usage: /curc/sw/install/bio/alphafold/2.3.1/alphafold-2.3.1/run\_alphaFold.sh <OPTIONS>

Required Parameters:

- d <data\_dir> Path to directory of supporting data
- o <output\_dir> Path to a directory that will store the results.
- f <fasta\_paths> Path to FASTA files containing sequences. If a FASTA file contains multiple sequences, then it will be folded as a multimer. To fold more sequences one after another, write the files separated by a comma
- t <max\_template\_date> Maximum template release date to consider (ISO-8601 format - i.e. YYYY-MM-DD). Important if folding historical test sets

To predict the structure of a protein in the PDB without using its experimental structure as a template, set **max\_template\_date** (-t) to before the release date of the structure.

TIP: A date like 1900-01-01 will ensure no templates are used.



# AlphaFold2 Optional Parameters

## Optional Parameters:

```
-g <use_gpu>          Enable NVIDIA runtime to run with GPUs (default: true)
-r <run_relax>         Whether to run the final relaxation step on the predicted models. Turning relax off might result in predictions with distracting stereochemical violations but might help in case you are having issues with the relaxation stage (default: true)
-e <enable_gpu_relax> Run relax on GPU if GPU is enabled (default: true)
-n <openmm_threads>    OpenMM threads (default: all available cores)
-a <gpu_devices>       Comma separated list of devices to pass to 'CUDA_VISIBLE_DEVICES' (default: 0)
-m <model_preset>      Choose preset model configuration - the monomer model, the monomer model with extra ensembling, monomer model with pTM head, or multimer model (default: 'monomer')
-c <db_preset>         Choose preset MSA database configuration - smaller genetic database config (reduced_dbs) or full genetic database config (full_dbs) (default: 'full_dbs')
```

Parameter values for --model\_preset (-m):

monomer

monomer\_casp14

monomer\_ptm

multimer



# AlphaFold2 Optional Parameters Cont'd

```
-p <use_precomputed_msas> Whether to read MSAs that have been written to disk. WARNING: This will not check if the sequence, database or configuration have changed (default: 'false')
-l <num_multimer_predictions_per_model> How many predictions (each with a different random seed) will be generated per model. E.g. if this is 2 and there are 5 models then there will be 10 predictions per input. Note: this FLAG only applies if model_preset=multimer (default: 5)
-b <benchmark> Run multiple JAX model evaluations to obtain a timing that excludes the compilation time, which should be more indicative of the time required for inferencing many proteins (default: 'false')
```

Using precomputed MSAs: AlphaFold looks for the MSAs in the output directory, so they must stay the same if you want to reuse them.



# Folding Monomers (1 Polypeptide Chain)

- The default model preset is monomer prediction (`--model_preset=monomer` or `-m "monomer"`)
- Requires a FASTA file with a single sequence

Run Command: `run_alphaFold -d $CURC_AF_DBs -o . -f $CURC_AF_EXAMPLES/dummy.fasta -t 2020-05-14 -m "monomer" -g true`



# Folding Multimers (>1 Polypeptide Chain)

- Requires a FASTA file containing **multiple** sequences

```
>bNACHT09
MKKIFISYSWDGEEHQEWVRKLADSLEEDHAYHVVWDGYDLDPLSDKNFFMERAATNTDFILVVATEKYKEKADNRNGGVGIETYLAARHWEMM
TGDKDKTNVITINREPDPSTPNYLNGHLVDFSKDELFNDSLVELKRLNAEPKFKRPEQKIGVTSLKSYNLTVKSDIIGIGAKNRICLINAEGT
DFSAGKKIKYEFWELRNPGGVITHVLALHNNIISQTLNRAAEIKERFQNISTLLLRPREKKGVESMDTILEKKGSRTVNINEYTYDEYIW
EYCIDSDFKNITPPDAIEFYTSQEVDSDKGIYSSGVNTLSDELLKTSDFYANLIIGGGIGKTSCLTVNKLIKEKSESFLTILRSEDIRKY
IHDAGINSLYVS DIYEMQAKYLNHNVNFDRNTFLLSILSGKIIIVDVGDELSSIFGDKFDLVFLKSLTKLHDELGETRILLTTRENNNSIS
QELISDLNINVYKLLGFKLDNCIKYLNIRFKHNEYKEMITSKIRSQIEKCSLSGEERIIPFFVDVISNIYEDNL YSGSEDDEFELSQUESTPYPSL
NEVNDHVIHSIFNREKVRHKYNISSEMIELFKTLNIELGDKWDINEASV LIGVLYENQQKELEDCILKNP LNNVSDSFTTYKYEFLHSYFNSL
SIMEFLTKDNVENTINLLSKSSKDTP EHKDIYKFLYNKGDALQILT PSLNKAIKMSLEYEKQDKIKYERAVAAIENIIYLMFTNGKKDEFTG
IIKKIYNSNQNDKITGLYIKGDLPPINLTGMAISHSKFKHYKRFLKSDFENS KFTY TQFINCHNENIKNTSFLAADIDRNTCDIGDLDYTFNMLT
KIQDAAEQVSEELAKFLSSFYKGSSFVENKKIYIKFSNRLPGLKGAFNKL SHGLISVKAKEVDTFYIISENYRKSVRRFLNDGYKDAKIKD
LIEFIKA
>T4_nrdC.3
MKTRKHYIDYFDSSLITKHDYQKGHREVINNIILRDFLDYIGWENHICKDTQNAYSHLSLLEWFKRSRLLSSVIAVNNVKKFMYPYIETNVSN
DNVVTFNIN DVKRTRYLEEFWSKDSKEKFASEFSHEFNNNNVMLFKHSRRLFCHGDDRTINVNVKDWTAKFIPSSQNGPFELLIIVCAPHEIYK
NLPYMKPC EANKHNTIRSLTYKLRTLLSKMDVVEFDDNTNYGLSLFETKVVIKLKD PNKFKP KPKPNHGNDTMKEEREYLSTRLIEVEKQIEE
HTKVLKDLTAKANGL RNAIEVLK
```

Run Command: `run_alphaFold -d $CURC_AF_DBs -o . -f $CURC_AF_EXAMPLES/multimer.fa -t 2020-05-14 -m “multimer” -g true`



# Folding Multimers (>1 Polypeptide Chain Cont'd)

```
I0521 14:26:53.124762 140296744294208 run_alphaFold.py:386] Have 25 models: ['model_1_multimer_v3_pred_0', 'model_1_multimer_v3_pred_1', 'model_1_multimer_v3_pred_2', 'model_1_multimer_v3_pred_3', 'model_1_multimer_v3_pred_4', 'model_2_multimer_v3_pred_0', 'model_2_multimer_v3_pred_1', 'model_2_multimer_v3_pred_2', 'model_2_multimer_v3_pred_3', 'model_2_multimer_v3_pred_4', 'model_3_multimer_v3_pred_0', 'model_3_multimer_v3_pred_1', 'model_3_multimer_v3_pred_2', 'model_3_multimer_v3_pred_3', 'model_3_multimer_v3_pred_4', 'model_4_multimer_v3_pred_0', 'model_4_multimer_v3_pred_1', 'model_4_multimer_v3_pred_2', 'model_4_multimer_v3_pred_3', 'model_4_multimer_v3_pred_4', 'model_5_multimer_v3_pred_0', 'model_5_multimer_v3_pred_1', 'model_5_multimer_v3_pred_2', 'model_5_multimer_v3_pred_3', 'model_5_multimer_v3_pred_4']
I0521 14:26:53.124969 140296744294208 run_alphaFold.py:403] Using random seed 21484416166174026
for the data pipeline
I0521 14:26:53.125209 140296744294208 run_alphaFold.py:161] Predicting multimer
I0521 14:26:53.139846 140296744294208 pipeline_multimer.py:210] Running monomer pipeline on chain A: bNACHT09
I0521 14:26:53.141395 140296744294208 jackhmmer.py:133] Launching subprocess "/curc/sw/install/bio/alphaFold/2.3.1/2.3.1_env/bin/jackhmmer -o /dev/null -A /scratch/alpine/lafr9499/tmpfqonmyqd/output.sto --noali --F1 0.0005 --F2 5e-05 --F3 5e-07 --incE 0.0001 -E 0.0001 --cpu 8 -N 1 /scratch/alpine/lafr9499/tmpic840x7q.fasta /gpfs/alpine1/datasets/bioinformatics/alphaFold/uniref90/uniref90.fasta"
```



# Example Job Script

- An example job script is located in  
`/curc/sw/install/bio/alphafold/2.3.1`
- Copy to any space you have write permissions and make your desired changes:

```
cd /projects/$USER
```

```
cp /curc/sw/install/bio/alphafold/2.3.1/alphafold_alphine.sh .
```



# Example Job Script

```
#!/bin/bash

#SBATCH --nodes=1
#SBATCH --time=06:00:00
#SBATCH --partition=aa100
#SBATCH --qos=normal
#SBATCH --gres=gpu:1
#SBATCH --job-name=monomer_test
#SBATCH --output=monomer_test_%j.out
#SBATCH --ntasks=8
#SBATCH --mail-type=ALL
#SBATCH --mail-user=<your email address>

module purge
module load alphafold

#change directory
cd /projects/$USER

#run AlphaFold
run_alphaFold -d $CURC_AF_DBs -o . -f $CURC_AF_EXAMPLES/dummy.fasta -t 2020-05-14 -m "monomer" -g true
```



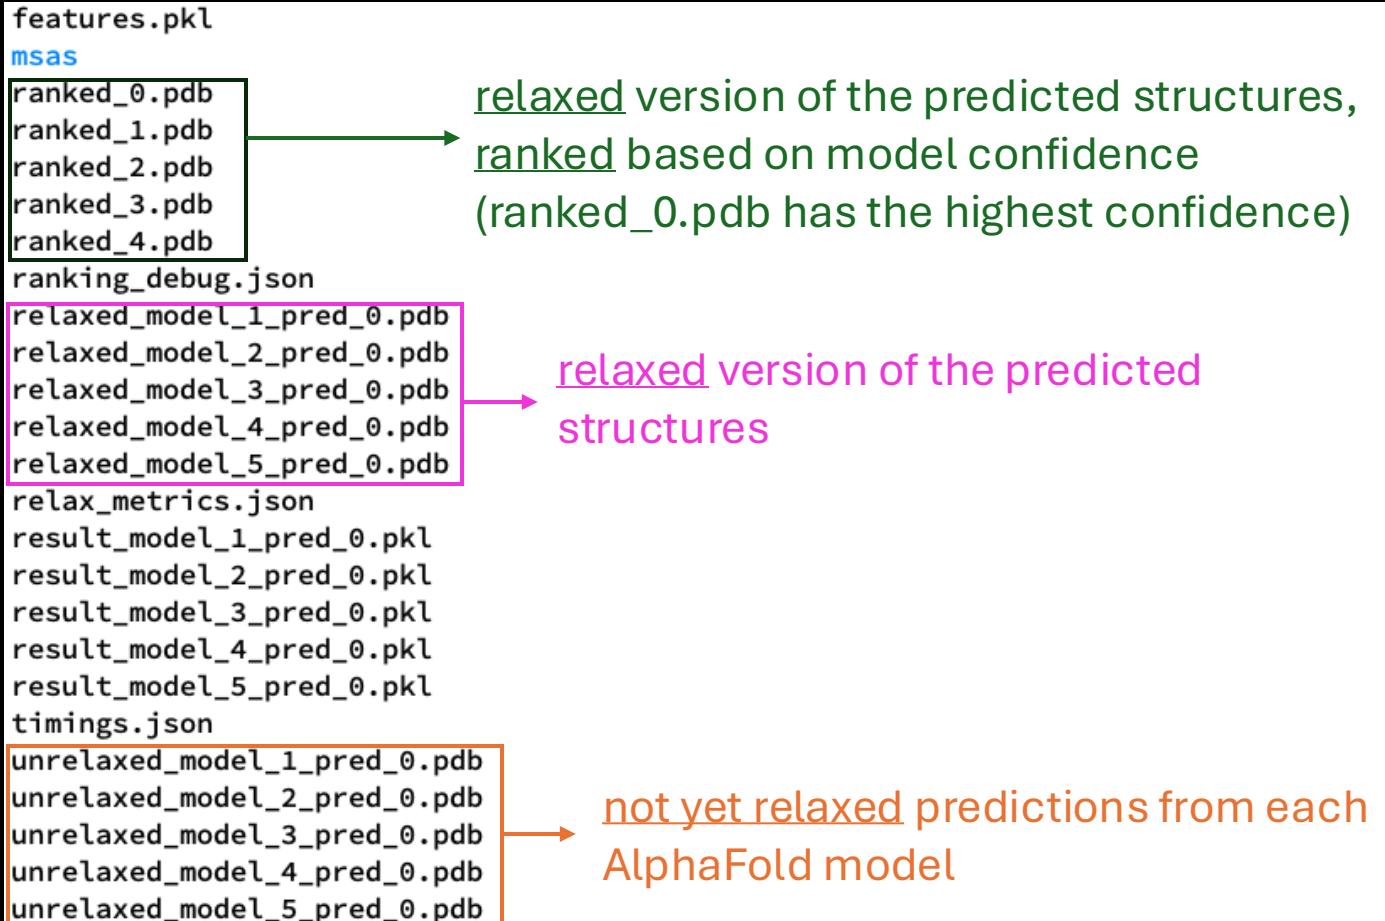
# AlphaFold2 Outputs Overview

- Slurm .out file (e.g.,  
`short_course_dummy_out_example.out`)
  - progress
  - warnings or errors
- a directory with the FASTA name (e.g., `dummy`)
  - computed MSAs
  - unrelaxed, relaxed, and ranked structures
  - raw model outputs
  - prediction metadata
  - timing information



# AlphaFold2 Outputs Explained 1

.pdb  
protein  
database  
format  
(human  
readable)



# AlphaFold2 Outputs Explained 2

.pkl  
pickle data  
format  
(binary; not  
readable)

```
features.pkl
msas
ranked_0.pdb
ranked_1.pdb
ranked_2.pdb
ranked_3.pdb
ranked_4.pdb
ranking_debug.json
relaxed_model_1_pred_0.pdb
relaxed_model_2_pred_0.pdb
relaxed_model_3_pred_0.pdb
relaxed_model_4_pred_0.pdb
relaxed_model_5_pred_0.pdb
relax_metrics.json
result_model_1_pred_0.pkl
result_model_2_pred_0.pkl
result_model_3_pred_0.pkl
result_model_4_pred_0.pkl
result_model_5_pred_0.pkl
timings.json
unrelaxed_model_1_pred_0.pdb
unrelaxed_model_2_pred_0.pdb
unrelaxed_model_3_pred_0.pdb
unrelaxed_model_4_pred_0.pdb
unrelaxed_model_5_pred_0.pdb
```

processed MSA information (NumPy arrays), used as input to produce structures

outputs in .pkl format. includes pLDDT and PTM values



# AlphaFold2 Outputs Explained 3

.json  
(readable  
data  
format)

```
features.pkl
msas
ranked_0.pdb
ranked_1.pdb
ranked_2.pdb
ranked_3.pdb
ranked_4.pdb
ranking_debug.json
relaxed_model_1_pred_0.pdb
relaxed_model_2_pred_0.pdb
relaxed_model_3_pred_0.pdb
relaxed_model_4_pred_0.pdb
relaxed_model_5_pred_0.pdb
relax_metrics.json
result_model_1_pred_0.pkl
result_model_2_pred_0.pkl
result_model_3_pred_0.pkl
result_model_4_pred_0.pkl
result_model_5_pred_0.pkl
timings.json
unrelaxed_model_1_pred_0.pdb
unrelaxed_model_2_pred_0.pdb
unrelaxed_model_3_pred_0.pdb
unrelaxed_model_4_pred_0.pdb
unrelaxed_model_5_pred_0.pdb
```

info on the pLDDT (confidence score) of each model and how they were ranked

relax metrics, including remaining violations

info on how long each section of the pipeline took to run



# AlphaFold2 Outputs Explained 4

Files in the  
msas  
directory  
are in  
various file  
formats, all  
human  
readable

```
features.pkl
msas
ranked_0.pdb
ranked_1.pdb
ranked_2.pdb
ranked_3.pdb
ranked_4.pdb
ranking_debug.json
relaxed_model_1_pred_0.pdb
relaxed_model_2_pred_0.pdb
relaxed_model_3_pred_0.pdb
relaxed_model_4_pred_0.pdb
relaxed_model_5_pred_0.pdb
relax_metrics.json
result_model_1_pred_0.pkl
result_model_2_pred_0.pkl
result_model_3_pred_0.pkl
result_model_4_pred_0.pkl
result_model_5_pred_0.pkl
timings.json
unrelaxed_model_1_pred_0.pdb
unrelaxed_model_2_pred_0.pdb
unrelaxed_model_3_pred_0.pdb
unrelaxed_model_4_pred_0.pdb
unrelaxed_model_5_pred_0.pdb
```

A directory containing the files describing  
the various genetic tool hits that were  
used to construct the input MSA



# Visualizing Outputs

- To date we have left this up to the user
- Some options are
  - AlphaPickle: <https://github.com/mattarnoldbio/alphapickle>
  - PyMol (free for education): <https://pymol.org/>
- AlphaFold Protein Structure Database website provides a great resource for learning how to interpret visualizations:  
<https://alphafold.ebi.ac.uk/entry/Q9Y223#help>



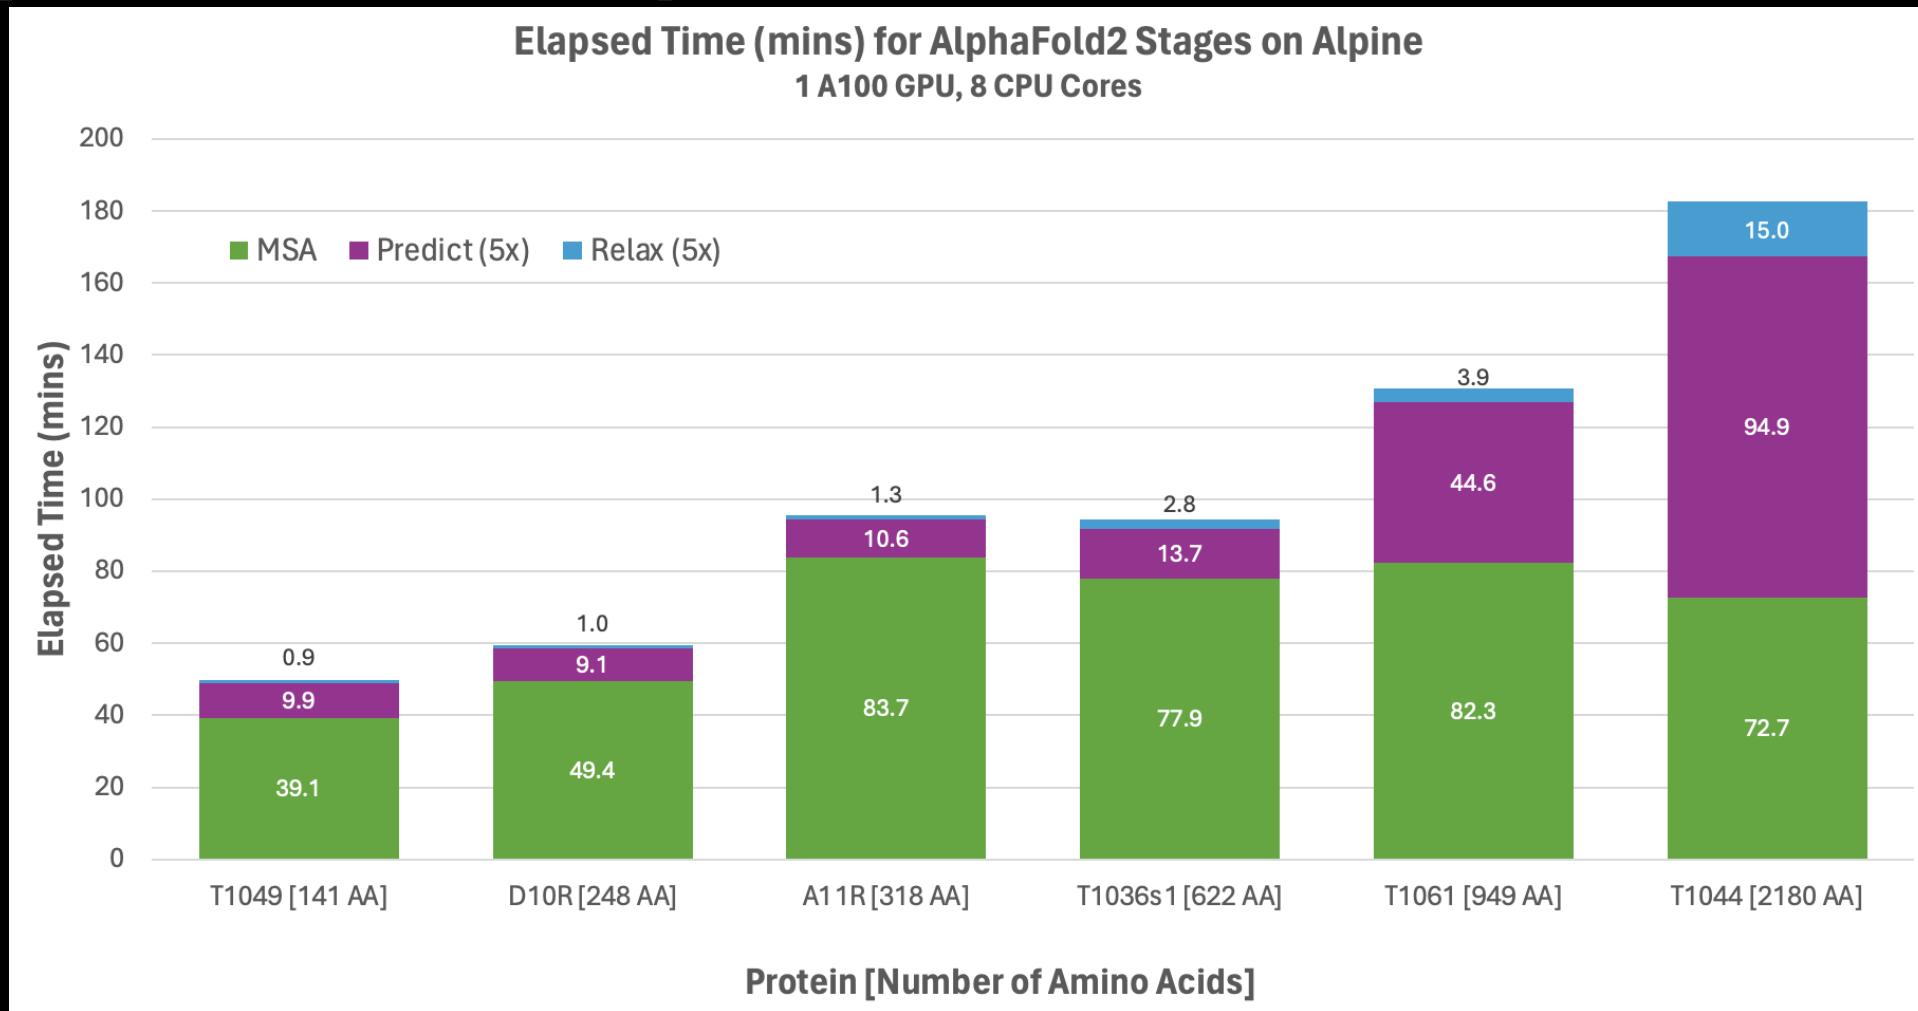
# AlphaFold2 Performance Considerations

```
ice set to true (defaults to false) to enable this.  
I0521 13:41:11.550761 140716613289792 run_alphafold.py:386] Have 5 models: ['model_1_pred_0', 'model_2_pred_0', 'mode  
l_3_pred_0', 'model_4_pred_0', 'model_5_pred_0']  
I0521 13:41:11.550974 140716613289792 run_alphafold.py:403] Using random seed 262200831417486581 for the data pipelin  
e  
I0521 13:41:11.551202 140716613289792 run_alphafold.py:161] Predicting dummy  
I0521 13:41:11.567026 140716613289792 jackhmmer.py:133] Launching subprocess "/curc/sw/install/bio/alphafold/2.3.1/2.  
3.1_env/bin/jackhmmer -o /dev/null -A /scratch/alpine/lafr9499/tmpuiv6sw6x/output.sto --noali --F1 0.0005 --F2 5e-05  
--F3 5e-07 --incE 0.0001 -E 0.0001 --cpu 8 -N 1 /curc/sw/install/bio/alphafold/examples/dummy.fasta /gpfs/alpine1/dat  
assets/bioinformatics/alphafold/uniref90/uniref90.fasta"  
I0521 13:41:11.584823 140716613289792 utils.py:36] Started Jackhmmer (uniref90.fasta) query
```

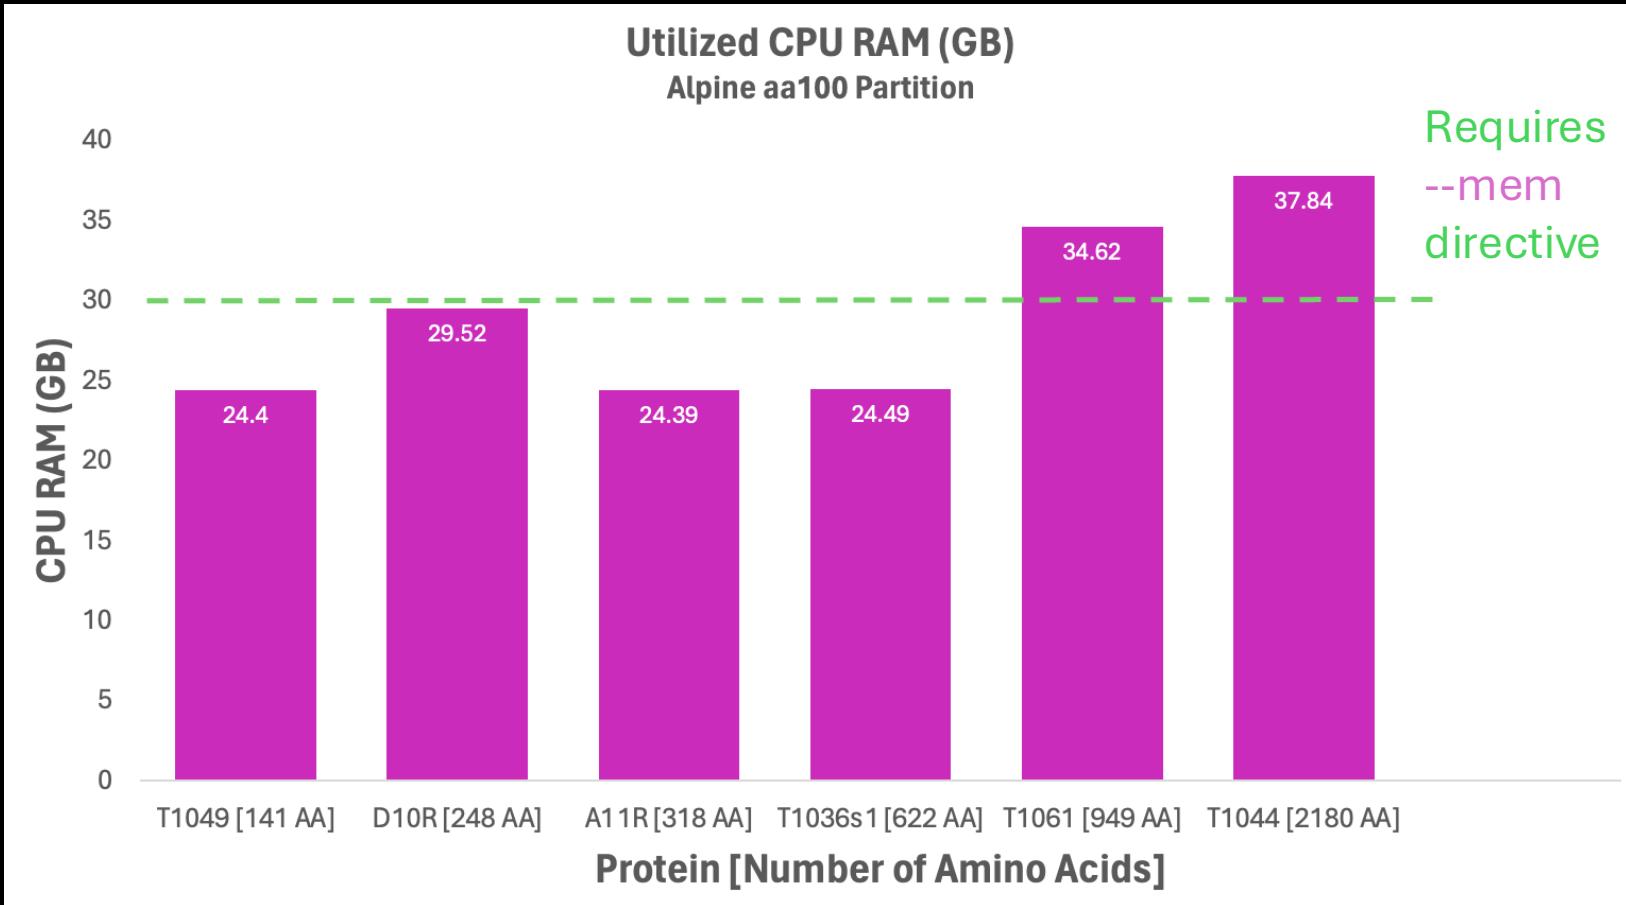
The number of cores/threads used during MSA steps is hardcoded and so AF cannot use more than 8 CPU cores. Requesting more cores will **not** improve performance and you may end up waiting longer for your job to start.



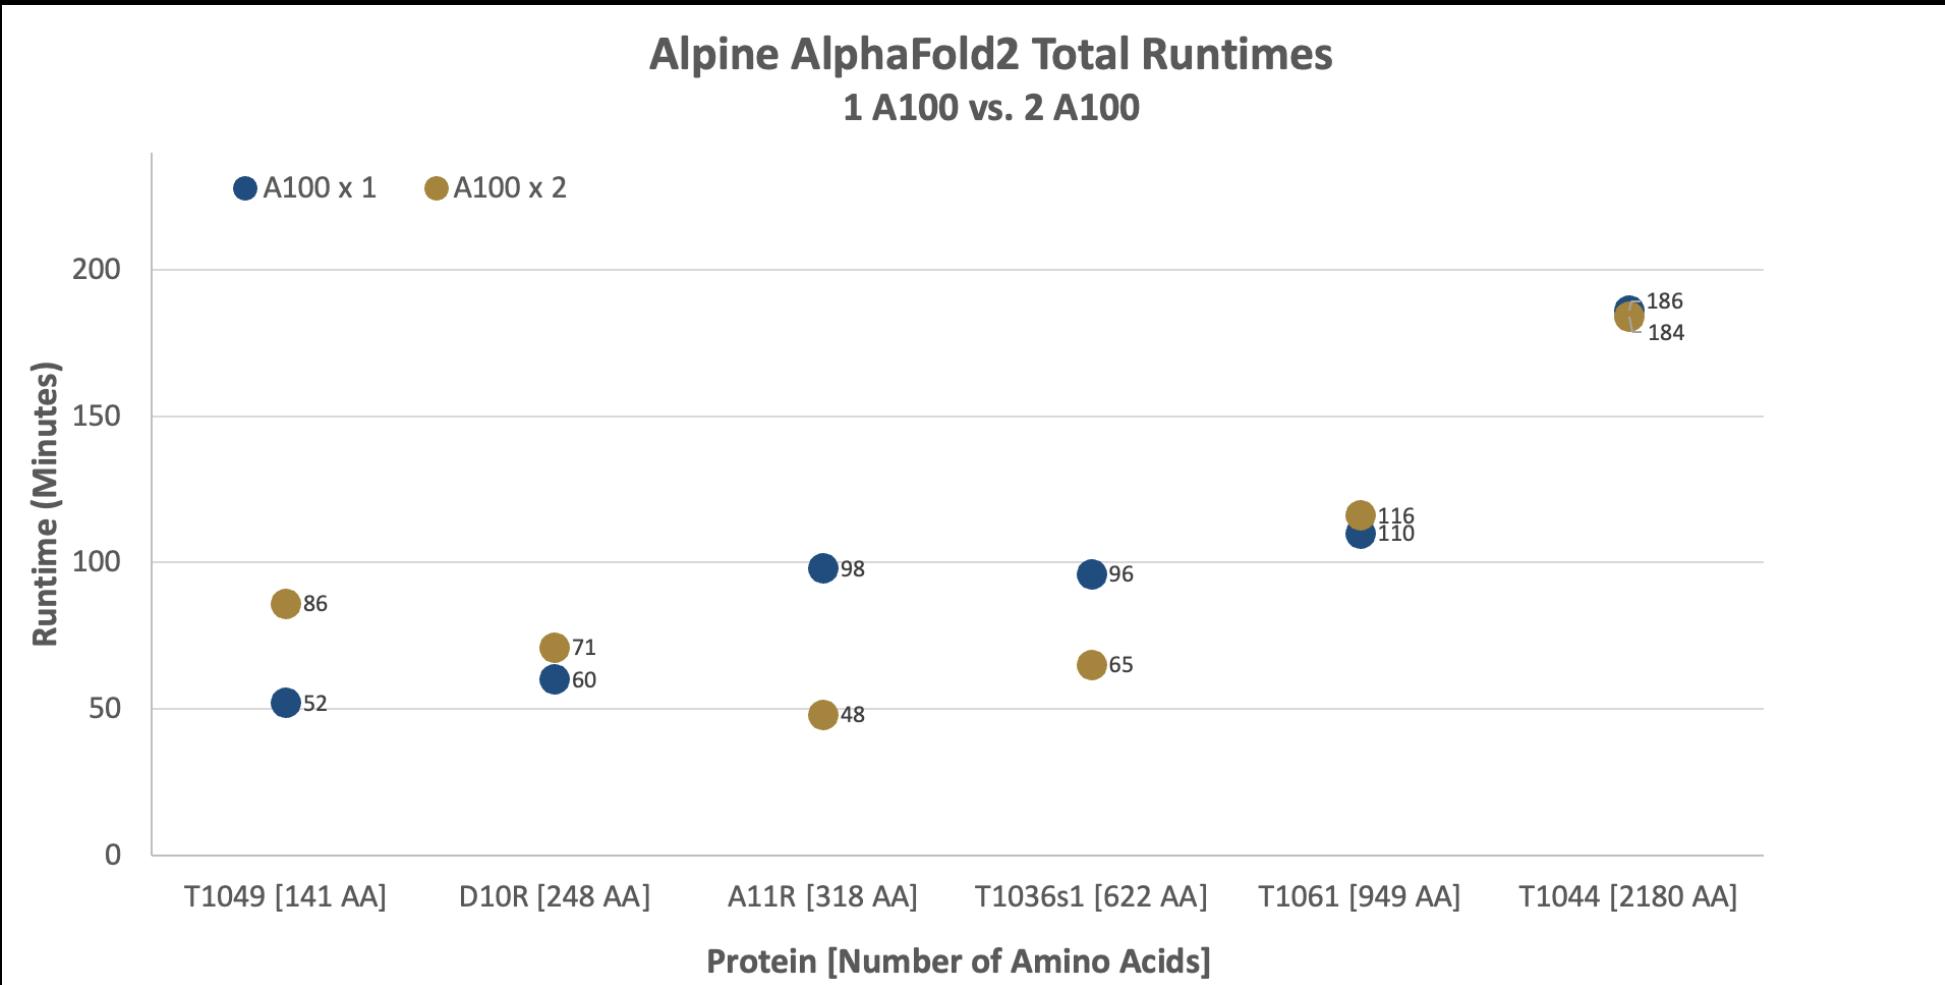
# AlphaFold2 Elapsed Time



# AlphaFold2 RAM Considerations



# AlphaFold2 Alpine Benchmarks



# AlphaFold Spin-Offs: ParaFold (ParallelFold)

- Divide CPU-intensive sections from the GPU-intensive sections
- PF isn't currently available as a module, but example scripts are available in `/curc/sw/install/bio/parafold/2.0`
  - `parafold_cpu.sh`
  - `parafold_gpu.sh`



# AlphaFold Spin-Offs: ParaFold Job Script

```
#!/bin/bash

#SBATCH --nodes=1
#SBATCH --time=10:00:00
#SBATCH --partition=amilan
#SBATCH --qos=normal
#SBATCH --job-name=parafold_cpu
#SBATCH --output=parafold_cpu_%j.out
#SBATCH --ntasks=8
#SBATCH --ntasks-per-node=8
#SBATCH --mem=60G
#SBATCH --mail-type=ALL
#SBATCH --mail-user=<your_email_address>

module purge
module load anaconda

conda activate /curc/sw/install/bio/parafold/2.0/2.0_env

export TMPDIR=/scratch/alpine/$USER

bash /curc/sw/install/bio/parafold/2.0/ParallelFold/run_alphaFold.sh -d /gpfs/alpine1/datasets/bioinformatics/alphafold
-o /projects/$USER/pf_runs -p multimer -i /curc/sw/install/bio/alphaFold/examples/multimer.fa -f -g false
```

-f tells PF to stop after generating the  
feature.pkl file



# AlphaFold Spin-Offs: ParaFold Job Script Cont'd

```
#!/bin/bash

#SBATCH --nodes=1
#SBATCH --time=24:00:00
#SBATCH --partition=aa100
#SBATCH --gres=gpu:1
#SBATCH --constraint=gpu80
#SBATCH --qos=normal
#SBATCH --job-name=parafold_gpu
#SBATCH --output=parafold_gpu_%j.out
#SBATCH --ntasks=8
#SBATCH --ntasks-per-node=8
#SBATCH --mail-type=ALL
#SBATCH --mail-user=<your_email_address>

module purge
module load cuda
module load anaconda

conda activate /curc/sw/install/bio/parafold/2.0/2.0_env

export TF_FORCE_UNIFIED_MEMORY=1
export XLA_PYTHON_CLIENT_MEM_FRACTION=4.0
export TMPDIR=/scratch/alpine/$USER

bash /curc/sw/install/bio/parafold/2.0/ParallelFold/run_alphaFold.sh -d /gpfs/alpine1/datasets/bioinformatics/alphafold
-o /projects/$USER/pf_runs -p multimer -i /curc/sw/install/bio/alphafold/examples/multimer.fa -g true -m model_1_multimer,model_2_multimer,model_3_multimer,model_4_multimer,model_5_multimer
```

Not needed for the multimer.fa example,  
but note that it is an option.

May need to add --mem directive for  
some predictions (relaxation step?)



# Explore AlphaFold2 on Alpine

Slides/Self-guided material:  
<https://bit.ly/af2alpine>



# Thank you

**Please give us your feedback!**

<http://tinyurl.com/curc-survey18>



Research Computing  
Office of Information and Technology

UNIVERSITY OF COLORADO BOULDER